

# Workshop in a compute cloud: not that obvious

Markus van Dijk  
SURFsara HPC Cloud advisor

# SURFsara

SARA was founded in 1971 by the two universities and the mathematical institute in Amsterdam for their computational needs.

Today, SURFsara is the “Dutch national high-performance computing and e-Science support center” and hosts the national supercomputer.

Academic research in the Netherlands can apply for free access to the resources.

# Systems at SURFsara

|                      | Cores | total RAM | RAM/core | storage  | GPUs |   |
|----------------------|-------|-----------|----------|----------|------|---|
| <b>Cartesius</b>     | 40960 | 117 TB    | 3 GB     | 7700 TB  | 132  | 1559 TFlops                                     |
| <b>Lisa</b>          | 8960  | 10 TB     | 1 GB     |          | -    | 158 TFlops                                      |
| <b>Grid</b>          | 15000 | ***       | 8 GB     | 26000 TB | -    |   |
| <b>Hadoop</b>        | 1576  | 10 TB     | 6 GB     | 2300 TB  | -    |   |
| <b>HPC Cloud</b>     | 1920  | 16 TB     | 8 GB     | 900 TB   | 20   |   |
| <b>Visualization</b> |       |           |          |          |      |   |
| <b>Archive</b>       |       |           |          | 25000 TB |      | tape  |
| <b>SURFdrive</b>     |       |           |          | 178 TB   |      | max 100GB/user<br>“dropbox” for<br>academic use |



# SURFsara HPC Cloud

Created for High Performance Computing:

- Fast private network between VMs (MPI).
- Large, fast disk storage (900TB Ceph).
- No overcommit, wait if full.

Stability and MTBF less important:

- No compute redundancy.
- No backups, but redundant storage.

No SLA :-)

# Demand for workshops

We had a growing demand for workshop support.

- By SURFsara: hands-on introduction to HPC Cloud as part of university curriculum.
- By institutes: as hands-on tooling training as part of their courses.

# Examples

- During 2015, 300 students from VU and AMC used 100.000 CoreHours during hands-on classes in bioinformatics and genomics.
- Visualization classes using Jupyter Notebooks in the cloud and Pandas, NetworkX, Folium (geological viz.).
- Hadoop training with Jupyter Notebooks in the Cloud using Spark to connect to SURFsara's Hadoop cluster.
- Hackathons

# Organization

The course organizer:

- plans and requests resources,
- prepares and tests the VM images,
- launches VMs, creates and distributes student logins,
- cleans up afterwards.

The cloud provider:

- makes sure the resources are available.

# Choices

- Students work all at the same time and in the same room, or spread out in time and space?
- Use 1 big VM for everybody or 1 each?
- Where are the individual results stored?
- What is done locally, what in the cloud?
- Do I need a plan B?



# What can go wrong?

- Local testing OK, cloud image not OK.
- Performance problems on scale up/out.
- A VM dies, data is lost.
- Resources (partially) unavailable.
- Cloud down or network connection failure.

# Your perspective

- How much can you trust the cloud provider to deliver the resources? What is your plan B?
- Do you have the skills to handle *many* VMs?
- Wouldn't you like the provider to handle VM startups and logins?
- Think about “high-availability” v.s. “prepare to fail”.
- What problem does Docker solve?

# Cloud provider perspective

- Capacity planning.
- Integration with Docker machine.
- Unpredictable usage by researchers → lock resources well in advance.
- SLA, Availability.
- High impact on current cloud, researchers suffer.

# Docker to the rescue?

## **Virtual Machine**

- “safe” environment
- full blown boot
- “normal” OS
- multi-user
- hard to test locally
- full install and maint.
- familiar technology

## **Docker container**

- not (yet) safe enough
- fast start/stop, small footprint
- single process
- multiple containers
- environment different
- reproducible build
- new (better!)

# VM meets Docker

- Use VMs to host Docker containers.
- Login and IP/port access management remains.
- Embrace Docker philosophy: prepare to fail.
- Need remains for simple supporting tools.
- VM and container technology will develop towards each other.

# Future at SURFsara

- SURFsara wants a separate cloud for non-research activities.
- Resource planing.
- Docker containers with predictable environment.
- Managing end-users:
  - Acceptance of terms of use.
  - Login names, passwords, public keys.
  - IP address/port for end-user access.
- WaaS: Workshop as a Service?
- Dynamic scaling?

# Thanks, any questions?

